

## Assignment : 11

**Name:** Shivraj Chaudar

**Div:** A **Class:** T.E

**Roll no. :** 3101018

### **#pom.xml**

```
<?xml version="1.0" encoding="UTF-8"?>
<project xmlns="http://maven.apache.org/POM/4.0.0"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/xsd/
maven-4.0.0.xsd">
  <modelVersion>4.0.0</modelVersion>

  <groupId>org.shivraj</groupId>
  <artifactId>WordCount</artifactId>
  <version>1.0-SNAPSHOT</version>

  <properties>
    <maven.compiler.source>8</maven.compiler.source>
    <maven.compiler.target>8</maven.compiler.target>
    <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
  </properties>

  <dependencies>
    <dependency>
      <groupId>org.apache.hadoop</groupId>
      <artifactId>hadoop-common</artifactId>
      <version>3.3.3</version>
    </dependency>
    <dependency>
      <groupId>org.apache.hadoop</groupId>
      <artifactId>hadoop-mapreduce-client-core</artifactId>
      <version>3.3.3</version>
    </dependency>
  </dependencies>
</project>
```

### **#WC\_Runner.java**

```
package org.shivraj;

import java.io.IOException;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapred.FileInputFormat;
```

```

import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.TextOutputFormat;
public class WC_Runner {
    public static void main(String[] args) throws IOException{
        JobConf conf = new JobConf(WC_Runner.class);
        conf.setJobName("WordCount");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(WC_Mapper.class);
        conf.setCombinerClass(WC_Reducer.class);
        conf.setReducerClass(WC_Reducer.class);
        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);
        FileInputFormat.setInputPaths(conf,new Path(args[0]));
        FileOutputFormat.setOutputPath(conf,new Path(args[1]));
        JobClient.runJob(conf);
    }
}

```

### ***#WC\_Mapper.java***

```

package org.shivraj;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

```

```

public class WC_Mapper extends MapReduceBase implements
Mapper<LongWritable,Text,Text,IntWritable>{

    private final static IntWritable one = new IntWritable(1);

    private Text word = new Text();

    public void map(LongWritable key, Text value,OutputCollector<Text,IntWritable>
output,Reporter reporter) throws IOException{

        String line = value.toString();

        StringTokenizer tokenizer = new StringTokenizer(line);

        while (tokenizer.hasMoreTokens()){

            word.set(tokenizer.nextToken());

            output.collect(word, one);

        }

    }
}

```

### ***#WC\_Reducer.java***

```

package org.shivraj;

import java.io.IOException;

import java.util.Iterator;

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapred.MapReduceBase;

import org.apache.hadoop.mapred.OutputCollector;

import org.apache.hadoop.mapred.Reducer;

import org.apache.hadoop.mapred.Reporter;

public class WC_Reducer extends MapReduceBase implements
Reducer<Text,IntWritable,Text,IntWritable> {

    public void reduce(Text key, Iterator<IntWritable> values,OutputCollector<Text,IntWritable>
output,Reporter reporter) throws IOException {

        int sum=0;

        while (values.hasNext()) {

            sum+=values.next().get();

        }

    }
}

```

```
output.collect(key,new IntWritable(sum));
```

```
}
```

```
}
```

### ***#Input.txt***

Hello wordcount MapReduce Hadoop program.

This is my first MapReduce program.

### **#Output**

**(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ %** hdfs dfs -mkdir -p /user/process

2025-03-31 12:25:41,823 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ %** hdfs dfs -put /Users/Shivrajchaudar/Desktop/input.txt /user/process

2025-03-31 12:26:49,178 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ %** hadoop jar /Users/Shivrajchaudar/hadoop-3.4.1/share/hadoop/hdfs/WordCount-1.0-SNAPSHOT.jar org.shivraj.WC\_Runner /user/process/input.txt /user/out

2025-03-31 12:35:46,244 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

2025-03-31 12:35:46,623 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /127.0.0.1:8032

2025-03-31 12:35:46,762 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /127.0.0.1:8032

2025-03-31 12:35:46,827 INFO mapreduce.JobSubmissionFiles: Permissions on staging directory /tmp/hadoop-yarn/staging/shivrajchaudar/.staging are incorrect: rwxrwxrwx. Fixing permissions to correct value rwx-----

2025-03-31 12:35:46,981 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.

2025-03-31 12:35:46,991 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/shivrajchaudar/.staging/job\_1743357063034\_0003

2025-03-31 12:35:47,160 INFO mapred.FileInputFormat: Total input files to process : 1

2025-03-31 12:35:48,025 INFO mapreduce.JobSubmitter: number of splits:2

2025-03-31 12:35:48,538 INFO mapreduce.JobSubmitter: Submitting tokens for job: job\_1743357063034\_0003

2025-03-31 12:35:48,538 INFO mapreduce.JobSubmitter: Executing with tokens: []

2025-03-31 12:35:48,675 INFO conf.Configuration: resource-types.xml not found

2025-03-31 12:35:48,675 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.

2025-03-31 12:35:48,779 INFO impl.YarnClientImpl: Submitted application application\_1743357063034\_0003

2025-03-31 12:35:48,799 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application\_1743357063034\_0003/

2025-03-31 12:35:48,799 INFO mapreduce.Job: Running job: job\_1743357063034\_0003

2025-03-31 12:35:53,962 INFO mapreduce.Job: Job job\_1743357063034\_0003 running in uber mode : false

2025-03-31 12:35:53,964 INFO mapreduce.Job: map 0% reduce 0%

2025-03-31 12:35:58,068 INFO mapreduce.Job: map 100% reduce 0%

2025-03-31 12:36:01,100 INFO mapreduce.Job: map 100% reduce 100%

2025-03-31 12:36:03,135 INFO mapreduce.Job: Job job\_1743357063034\_0003 completed successfully

2025-03-31 12:36:03,248 INFO mapreduce.Job: Counters: 50

#### File System Counters

FILE: Number of bytes read=150

FILE: Number of bytes written=928700

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=311

HDFS: Number of bytes written=77

HDFS: Number of read operations=11

HDFS: Number of large read operations=0

HDFS: Number of write operations=2

HDFS: Number of bytes read erasure-coded=0

#### Job Counters

Launched map tasks=2

Launched reduce tasks=1

Data-local map tasks=2

Total time spent by all maps in occupied slots (ms)=6600

Total time spent by all reduces in occupied slots (ms)=3034

Total time spent by all map tasks (ms)=3300

Total time spent by all reduce tasks (ms)=1517

Total vcore-milliseconds taken by all map tasks=3300

Total vcore-milliseconds taken by all reduce tasks=1517

Total megabyte-milliseconds taken by all map tasks=3379200

Total megabyte-milliseconds taken by all reduce tasks=1553408

#### Map-Reduce Framework

Map input records=2

Map output records=11

Map output bytes=122

Map output materialized bytes=156

Input split bytes=192

Combine input records=11

Combine output records=11

Reduce input groups=9

Reduce shuffle bytes=156

Reduce input records=11

Reduce output records=9

Spilled Records=22

Shuffled Maps =2

Failed Shuffles=0

Merged Map outputs=2

GC time elapsed (ms)=96  
CPU time spent (ms)=0  
Physical memory (bytes) snapshot=0  
Virtual memory (bytes) snapshot=0  
Total committed heap usage (bytes)=724041728  
Shuffle Errors  
BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0  
WRONG\_REDUCE=0  
File Input Format Counters  
Bytes Read=119  
File Output Format Counters  
Bytes Written=77

**(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ % hdfs dfs -ls /user/out**

2025-03-31 12:36:28,522 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

-rw-r--r--	3	shivrajchaudar	sharedgroup	0	2025-03-31 12:35	/user/out/_SUCCESS
-rw-r--r--	3	shivrajchaudar	sharedgroup	77	2025-03-31 12:35	/user/out/part-00000

**(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ % hdfs dfs -cat /user/out/part-00000**

2025-03-31 12:37:04,184 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Hadoop 1

Hello 1

MapReduce 2

This 1

first 1

is 1

my 1

program. 2

wordcount 1

**(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ %**

```

(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ % hdfs dfs -mkdir -p /user/process
2025-03-31 12:25:41,823 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ % hdfs dfs -put /Users/Shivrajchaudar/Desktop/input.txt /user/process
2025-03-31 12:26:49,178 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ % hadoop jar /Users/Shivrajchaudar/hadoop-3.4.1/share/hadoop/hdfs/WordCount-1.0-SNAPSHOT.jar org.shivraj.WC_Runner /user/process/input.txt /user/out
2025-03-31 12:35:46,244 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2025-03-31 12:35:46,623 INFO client.DefaultNotHARMAFailoverProxyProvider: Connecting to ResourceManager at /127.0.0.1:8032
2025-03-31 12:35:46,762 INFO client.DefaultNotHARMAFailoverProxyProvider: Connecting to ResourceManager at /127.0.0.1:8032
2025-03-31 12:35:46,827 INFO mapreduce.JobSubmissionFiles: Permissions on staging directory /tmp/hadoop-yarn/staging/shivrajchaudar/.staging are incorrect: rwxrwxrwx. Fixing permissions to correct value rwxrwxrwx.
2025-03-31 12:35:46,981 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-03-31 12:35:46,991 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/shivrajchaudar/.staging/job_1743357063034_0003
2025-03-31 12:35:47,160 INFO mapred.FileInputFormat: Total input files to process : 1
2025-03-31 12:35:48,025 INFO mapreduce.JobSubmitter: number of splits:2
2025-03-31 12:35:48,538 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1743357063034_0003
2025-03-31 12:35:48,538 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-03-31 12:35:48,675 INFO conf.Configuration: resource-types.xml not found
2025-03-31 12:35:48,675 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2025-03-31 12:35:48,779 INFO impl.YarnClientImpl: Submitted application application_1743357063034_0003
2025-03-31 12:35:48,799 INFO mapreduce.Job: The url to track the job: http://localhost:8080/proxy/application_1743357063034_0003/
2025-03-31 12:35:48,799 INFO mapreduce.Job: Running job: job_1743357063034_0003
2025-03-31 12:35:53,962 INFO mapreduce.Job: Job job_1743357063034_0003 running in uber mode : false
2025-03-31 12:35:53,964 INFO mapreduce.Job: map 0% reduce 0%
2025-03-31 12:35:58,068 INFO mapreduce.Job: map 100% reduce 0%
2025-03-31 12:36:01,100 INFO mapreduce.Job: map 100% reduce 100%
2025-03-31 12:36:03,135 INFO mapreduce.Job: Job job_1743357063034_0003 completed successfully
2025-03-31 12:36:03,248 INFO mapreduce.Job: Counters: 50
File System Counters
  FILE: Number of bytes read=150
  FILE: Number of bytes written=928700
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=311
  HDFS: Number of bytes written=77
  HDFS: Number of read operations=11
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=6600
  Total time spent by all reduces in occupied slots (ms)=3034
  Total time spent by all map tasks (ms)=3300
  Total time spent by all reduce tasks (ms)=1517
  Total vcore-milliseconds taken by all map tasks=3300
  Total vcore-milliseconds taken by all reduce tasks=1517
  Total megabyte-milliseconds taken by all map tasks=3379200
  Total megabyte-milliseconds taken by all reduce tasks=1553400
Map-Reduce Framework
  Map input records=2
  Map output records=11
  Map output bytes=122
  Map output materialized bytes=156
  Input split bytes=192
  Combine input records=11
  Combine output records=11

```

```

  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=6600
  Total time spent by all reduces in occupied slots (ms)=3034
  Total time spent by all map tasks (ms)=3300
  Total time spent by all reduce tasks (ms)=1517
  Total vcore-milliseconds taken by all map tasks=3300
  Total vcore-milliseconds taken by all reduce tasks=1517
  Total megabyte-milliseconds taken by all map tasks=3379200
  Total megabyte-milliseconds taken by all reduce tasks=1553400
Map-Reduce Framework
  Map input records=2
  Map output records=11
  Map output bytes=122
  Map output materialized bytes=156
  Input split bytes=192
  Combine input records=11
  Combine output records=11
  Reduce input groups=9
  Reduce shuffle bytes=156
  Reduce input records=11
  Reduce output records=9
  Spilled Records=22
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=96
  CPU time spent (ms)=0
  Physical memory (bytes) snapshot=0
  Virtual memory (bytes) snapshot=0
  Total committed heap usage (bytes)=724041728
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=119
File Output Format Counters
  Bytes Written=77
(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ % hdfs dfs -ls /user/out
2025-03-31 12:36:28,522 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 3 shivrajchaudar sharedgroup 0 2025-03-31 12:35 /user/out/_SUCCESS
-rw-r--r-- 3 shivrajchaudar sharedgroup 77 2025-03-31 12:35 /user/out/part-00000
(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ % hdfs dfs -cat /user/out/part-00000
2025-03-31 12:37:04,184 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Hadoop 1
Hello 1
MapReduce 2
This 1
first 1
is 1
my 1
program. 2
wordcount 1
(base) shivrajchaudar@Shivrajs-Macbook-Pro ~ %

```

Browsing HDFS

localhost:9870/explorer.html#/user/process

Home | Microsoft...ChatGPTGmailYouTubeMapsWelcome – Google...leetcodepwskills.com/learn/Online Shopping si...Home | TraveleaseAll Bookmarks

HadoopOverviewDatanodesDatanode Volume FailuresSnapshotStartup ProgressUtilities

# Browse Directory

Go!

Show 25 entries

Search:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	<a href="#">-rw-r--r--</a>	<a href="#">shivrajchaudar</a>	<a href="#">sharedgroup</a>	79 B	Mar 31 12:26	<a href="#">3</a>	128 MB	<a href="#">input.txt</a>	<div></div>

Showing 1 to 1 of 1 entries

Previous1Next

Hadoop, 2024.

Browsing HDFS

localhost:9870/explorer.html#/user/out

Home | Microsoft...ChatGPTGmailYouTubeMapsWelcome – Google...leetcodepwskills.com/learn/Online Shopping si...Home | TraveleaseAll Bookmarks

HadoopOverviewDatanodesDatanode Volume FailuresSnapshotStartup ProgressUtilities

# Browse Directory

Go!

Show 25 entries

Search:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	<a href="#">-rw-r--r--</a>	<a href="#">shivrajchaudar</a>	<a href="#">sharedgroup</a>	0 B	Mar 31 12:35	<a href="#">3</a>	128 MB	<a href="#">_SUCCESS</a>	<div></div>
<input type="checkbox"/>	<a href="#">-rw-r--r--</a>	<a href="#">shivrajchaudar</a>	<a href="#">sharedgroup</a>	77 B	Mar 31 12:35	<a href="#">3</a>	128 MB	<a href="#">part-00000</a>	<div></div>

Showing 1 to 2 of 2 entries

Previous1Next

Hadoop, 2024.



Browsing HDFS

localhost:9870/explorer.html#/user/out

Home | Microsoft...ChatGPTGmailYouTubeMapsWelcome – Google...leetcodepwwskills.com/learn/Online Shopping si...Home | Travellease

All Bookmarks

HadoopOverviewDatanodes

# Browse Directory

/user/out

Show 25 entries

<input type="checkbox"/>	Permission	Owner
<input type="checkbox"/>	-rw-r--r--	shivrajch
<input type="checkbox"/>	-rw-r--r--	shivrajch

Showing 1 to 2 of 2 entries

Hadoop, 2024.

File information - part-00000

[Download](#)[Head the file \(first 32K\)](#)[Tail the file \(last 32K\)](#)

Block information -- Block 0

Block ID: 1073741857

Block Pool ID: BP-2095407043-127.0.0.1-1743357293385

Generation Stamp: 1033

Size: 77

Availability:

- localhost

File contents

Hadoop 1

Hello 1

MapReduce 2

This 1

first 1

is 1

my 1

program. 2

wordcount 1

Close